


(a) per-pixel feature vector



one row per labelled pixel; label from the polygon's 'Id' property

(b) split and cross-validation are grouped by polygon

Pixels inside one training polygon are near-duplicates. A pixel-level split puts copies of the same observation on both sides of the divide, so the hold-out score measures memory rather than generalisation. Every split here - the 70/30 hold-out and the 5 tuning folds - keeps a polygon whole.



(c) classifiers compared



identical features, identical split, identical tuning budget - the only difference between the six maps is the decision rule